

# Evaluating Performance Measures

[Code ▼](#)

The objective of this assignment is to compare two different classification algorithms using accuracy and performance metrics. I'm going to use K-NN and Naive Bayes to predict the age of an abalone using abalone features. Abalone are shell fish that are popular to eat in many countries, especially raw in a sashimi spread. The rings attribute corresponds to an abalone's age in years (after the abalone reaches 1 to 1.5 years of age). The process of determining an abalone's age is tedious and time consuming, so using classification machine learning might be useful for predicting an abalone's age.

Data was taken from this website <https://archive.ics.uci.edu/ml/datasets/Abalone> (<https://archive.ics.uci.edu/ml/datasets/Abalone>)

Here are the attribute descriptions:

Sex / nominal / – / M, F, and I (infant) Length / continuous / mm / Longest shell measurement Diameter / continuous / mm / perpendicular to length Height / continuous / mm / with meat in shell Whole weight / continuous / grams / whole abalone Shucked weight / continuous / grams / weight of meat Viscera weight / continuous / grams / gut weight (after bleeding) Shell weight / continuous / grams / after being dried Rings / integer / – / +1.5 gives the age in years

[Hide](#)[Hide](#)

```
abalone <- read.csv(url("https://archive.ics.uci.edu/ml/machine-learning-databases/a
balone/abalone.data"), header = FALSE, sep = ",")
colnames(abalone) <- c("sex", "length", 'diameter', 'height', 'whole_weight', 'shucked_weight', 'viscera_weight', 'shell_weight', 'rings' )
```

[Hide](#)[Hide](#)

```
summary(abalone)
```

sex	length	diameter	height
F:1307	Min. :0.075	Min. :0.0550	Min. :0.0000
I:1342	1st Qu.:0.450	1st Qu.:0.3500	1st Qu.:0.1150
M:1528	Median :0.545	Median :0.4250	Median :0.1400
	Mean :0.524	Mean :0.4079	Mean :0.1395
	3rd Qu.:0.615	3rd Qu.:0.4800	3rd Qu.:0.1650
	Max. :0.815	Max. :0.6500	Max. :1.1300

  

whole_weight	shucked_wieght	viscera_wieght
Min. :0.0020	Min. :0.0010	Min. :0.0005
1st Qu.:0.4415	1st Qu.:0.1860	1st Qu.:0.0935
Median :0.7995	Median :0.3360	Median :0.1710
Mean :0.8287	Mean :0.3594	Mean :0.1806
3rd Qu.:1.1530	3rd Qu.:0.5020	3rd Qu.:0.2530
Max. :2.8255	Max. :1.4880	Max. :0.7600

  

shell_weight	rings
Min. :0.0015	Min. : 1.000
1st Qu.:0.1300	1st Qu.: 8.000
Median :0.2340	Median : 9.000
Mean :0.2388	Mean : 9.934
3rd Qu.:0.3290	3rd Qu.:11.000
Max. :1.0050	Max. :29.000

[Hide](#)
[Hide](#)

```
str(abalone)
```

```
'data.frame': 4177 obs. of 9 variables:
 $ sex      : Factor w/ 3 levels "F","I","M": 3 3 1 3 2 2 1 1 3 1 ...
 $ length   : num  0.455 0.35 0.53 0.44 0.33 0.425 0.53 0.545 0.475 0.55 ...
 $ diameter : num  0.365 0.265 0.42 0.365 0.255 0.3 0.415 0.425 0.37 0.44 ...
 $ height   : num  0.095 0.09 0.135 0.125 0.08 0.095 0.15 0.125 0.125 0.15 ...
 $ whole_weight : num  0.514 0.226 0.677 0.516 0.205 ...
 $ shucked_wieght: num  0.2245 0.0995 0.2565 0.2155 0.0895 ...
 $ viscera_wieght: num  0.101 0.0485 0.1415 0.114 0.0395 ...
 $ shell_weight : num  0.15 0.07 0.21 0.155 0.055 0.12 0.33 0.26 0.165 0.32 ...
 $ rings      : int  15 7 9 10 7 8 20 16 9 19 ...
```

[Hide](#)
[Hide](#)

```
summary(abalone$rings)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
1.000	8.000	9.000	9.934	11.000	29.000

As shown above, the “rings” variable has a range between 1-29. This is the variable that we want to predict, and predicting this many levels might not give us the insight we’re looking for. I suspect that there’s an optimal age range for harvesting abalones for consumption. While I don’t know this age range, this project could be adjusted with the sought-after age range inserted. For now, we’ll break the rings variable into 3 levels “young” for abalones less than 8, “adult” for abalones between 8-11, and “old” for abalones older than 11.

Hide

Hide

```
abalone$rings <- as.numeric(abalone$rings)
abalone$rings <- cut(abalone$rings, br=c(-1,8,11,35), labels = c("young", 'adult', 'old'))
abalone$rings <- as.factor(abalone$rings)
summary(abalone$rings)
```

young	adult	old
1407	1810	960

I’m going to create a couple of different classification models, and then compare them using accuracy and performance metrics. I’ll start with a KNN classification algorithm. Because KNN requires all numeric variables for prediction, I’m going to remove the “sex” variable.

Hide

Hide

```
z <- abalone
z$sex <- NULL
```

I’ll now normalize the data using min max normalization

Hide

Hide

```
normalize <- function(x) {
  return ((x - min(x)) / (max(x) - min(x)))
}
z[1:7] <- as.data.frame(lapply(z[1:7], normalize))
summary(z$shucked_weight)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0000	0.1244	0.2253	0.2410	0.3369	1.0000

Now each variable has a min of 0 and a max of 1. We’ll now split the data into training and testing sets.

Hide

Hide

```
ind <- sample(2, nrow(z), replace=TRUE, prob=c(0.7, 0.3))
KNNtrain <- z[ind==1,]
KNNtest <- z[ind==2,]
```

Now we run the model. I'm going to make k equal to the square root of 2918, the number of observations in the training set.

Hide

Hide

```
library(class)
KNNpred <- knn(train = KNNtrain[1:7], test = KNNtest[1:7], cl = KNNtrain$strings, k = 5
4)
```

Let's see how the model does on the test data.

Hide

Hide

```
library("gmodels")
CrossTable(x = KNNtest$strings, y = KNNpred, prop.chisq = FALSE)
```

## Cell Contents

N
N / Row Total
N / Col Total
N / Table Total

Total Observations in Table: 1276

	KNNpred			
KNNtest\$strings	young	adult	old	Row Total
-----				
young	325	93	0	418
	0.778	0.222	0.000	0.328
	0.747	0.135	0.000	
	0.255	0.073	0.000	
-----				
adult	90	418	35	543
	0.166	0.770	0.064	0.426
	0.207	0.607	0.230	
	0.071	0.328	0.027	
-----				
old	20	178	117	315
	0.063	0.565	0.371	0.247
	0.046	0.258	0.770	
	0.016	0.139	0.092	
-----				
Column Total	435	689	152	1276
	0.341	0.540	0.119	
-----				

Hide

Hide

```
(328+451+97)/((84+2+26+91+21+159)+(328+451+97))
```

```
[1] 0.6957903
```

This KNN classifier predicted the abalone age with 69% accuracy - likely not accurate enough for an abalone harvester to trust. Before moving on to more specific accuracy and performance tests I'm going to try a smaller k value and see if it improves the accuracy.

[Hide](#)[Hide](#)

```
library(class)
KNNpred <- knn(train = KNNtrain[1:7], test = KNNtest[1:7], cl = KNNtrain$strings, k = 1
0)
```

[Hide](#)[Hide](#)

```
library("gmodels")
CrossTable(x = KNNtest$strings, y = KNNpred, prop.chisq = FALSE)
```

## Cell Contents

N
N / Row Total
N / Col Total
N / Table Total

Total Observations in Table: 1276

	KNNpred			
KNNtest\$strings	young	adult	old	Row Total
young	318	96	4	418
	0.761	0.230	0.010	0.328
	0.741	0.150	0.019	
	0.249	0.075	0.003	
adult	94	382	67	543
	0.173	0.703	0.123	0.426
	0.219	0.597	0.324	
	0.074	0.299	0.053	
old	17	162	136	315
	0.054	0.514	0.432	0.247
	0.040	0.253	0.657	
	0.013	0.127	0.107	
Column Total	429	640	207	1276
	0.336	0.502	0.162	

Hide

Hide

$(313+422+135)/((313+422+135)+(94+7+57+89+17+125))$

[1] 0.6910246

This model has just about the same predictive power on the test set. This can also be shown in the confusion matrix below:

Hide

[Hide](#)

```
library(caret)
```

```
package 'caret' was built under R version 3.3.2Loading required package: lattice
Loading required package: ggplot2
package 'ggplot2' was built under R version 3.3.2
```

[Hide](#)[Hide](#)

```
confusionMatrix(KNNpred, KNNtest$strings)
```

### Confusion Matrix and Statistics

```

      Reference
Prediction young adult old
   young    318     94  17
   adult     96    382 162
   old         4     67 136
```

### Overall Statistics

```

      Accuracy : 0.6552
      95% CI   : (0.6284, 0.6813)
No Information Rate : 0.4255
P-Value [Acc > NIR] : < 2.2e-16
```

```

      Kappa : 0.4581
McNemar's Test P-Value : 2.749e-10
```

### Statistics by Class:

	Class: young	Class: adult	Class: old
Sensitivity	0.7608	0.7035	0.4317
Specificity	0.8706	0.6480	0.9261
Pos Pred Value	0.7413	0.5969	0.6570
Neg Pred Value	0.8819	0.7469	0.8326
Prevalence	0.3276	0.4255	0.2469
Detection Rate	0.2492	0.2994	0.1066
Detection Prevalence	0.3362	0.5016	0.1622
Balanced Accuracy	0.8157	0.6758	0.6789

The misclassification rate is 1 minus the accuracy, shown below.

[Hide](#)[Hide](#)



```
1-0.691
```

```
[1] 0.309
```

Let's now create a naive bayes classifier for the same data.

[Hide](#)[Hide](#)

```
NBtrain <- KNNtrain  
NBtest <- KNNtest
```

[Hide](#)[Hide](#)

```
library(e1071)
```

```
package 'e1071' was built under R version 3.3.2
```

[Hide](#)[Hide](#)

```
model <- naiveBayes(rings ~., data = NBtrain)  
model
```

Naive Bayes Classifier for Discrete Predictors

Call:  
naiveBayes.default(x = X, y = Y, laplace = laplace)

A-priori probabilities:

```
Y  
      young      adult      old  
0.3409169 0.4367459 0.2223371
```

Conditional probabilities:

```
      length  
Y      [,1]      [,2]  
young 0.4716203 0.1490067  
adult 0.6712654 0.1177351  
old    0.6912424 0.1092433
```

```
      diameter  
Y      [,1]      [,2]  
young 0.4507906 0.1502746
```

```
adult 0.6580688 0.1193139
old    0.6863657 0.1123703
```

height

```
Y          [,1]      [,2]
young 0.09493365 0.04008618
adult  0.13448953 0.02663204
old    0.14618234 0.02554316
```

whole\_weight

```
Y          [,1]      [,2]
young 0.1544802 0.1083561
adult  0.3494055 0.1509557
old    0.3924864 0.1602648
```

shucked\_wieght

```
Y          [,1]      [,2]
young 0.1337806 0.09876195
adult  0.2970106 0.14060814
old    0.2962226 0.13836111
```

viscera\_wieght

```
Y          [,1]      [,2]
young 0.1241098 0.09026025
adult  0.2866922 0.12776408
old    0.3145439 0.13646942
```

shell\_weight

```
Y          [,1]      [,2]
young 0.1206878 0.08019425
adult  0.2733141 0.10873391
old    0.3365036 0.13167884
```

Hide

Hide

```
pred <- predict(model, NBtest)
print(confusionMatrix(pred,NBtest$rings))
```

Confusion Matrix and Statistics

Reference				
Prediction	young	adult	old	
young	336	126	47	
adult	76	285	139	
old	6	132	129	

Overall Statistics

Accuracy : 0.5878  
95% CI : (0.5602, 0.6149)  
No Information Rate : 0.4255  
P-Value [Acc > NIR] : < 2.2e-16  
  
Kappa : 0.3667  
Mcnemar's Test P-Value : 1.32e-09

Statistics by Class:

	Class: young	Class: adult	Class: old
Sensitivity	0.8038	0.5249	0.4095
Specificity	0.7984	0.7067	0.8564
Pos Pred Value	0.6601	0.5700	0.4831
Neg Pred Value	0.8931	0.6675	0.8157
Prevalence	0.3276	0.4255	0.2469
Detection Rate	0.2633	0.2234	0.1011
Detection Prevalence	0.3989	0.3918	0.2092
Balanced Accuracy	0.8011	0.6158	0.6330

The accuracy rate for the naive bayes model predicting the test set is only about 59%, which makes the misclassification rate approx. 41%.

While it’s likely that neither algorithm is adequate for predicting the abalone age, the KNN model is more accurate so far.

Let’s try a bootstrapping method for further model evaluation.

Hide

Hide

```
library(caret)
train_control <- trainControl(method='boot', number = 100)

trModel <- train(rings~., data = z, trControl=train_control, method="nb")
```

Hide

Hide

```
print(trModel)
```

Naive Bayes

```
4177 samples
  7 predictor
  3 classes: 'young', 'adult', 'old'
```

No pre-processing

Resampling: Bootstrapped (100 reps)

Summary of sample sizes: 4177, 4177, 4177, 4177, 4177, 4177, ...

Resampling results across tuning parameters:

usekernel	Accuracy	Kappa
FALSE	0.5810706	0.3529160
TRUE	0.6069231	0.3799894

Tuning parameter 'fL' was held constant at a value of 0

Tuning parameter 'adjust' was held constant at a value of 1

Accuracy was used to select the optimal model using the  
largest value.

The final values used for the model were fL = 0, usekernel =  
TRUE and adjust = 1.

Hide

Hide

```
trModel2 <- train(rings~., data = z, trControl=train_control, method="knn")
```

Hide

Hide

```
print(trModel2)
```

## k-Nearest Neighbors

```
4177 samples
  7 predictor
  3 classes: 'young', 'adult', 'old'
```

No pre-processing

Resampling: Bootstrapped (100 reps)

Summary of sample sizes: 4177, 4177, 4177, 4177, 4177, 4177, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.6222626	0.4091962
7	0.6347705	0.4266619
9	0.6447592	0.4406755

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was k = 9.

The bootstrapping method indicates that the KNN model might be slightly more accurate for classifying the data. It was also estimated that the most effective k value for the KNN model would be 9. We used 10 for our model.

Let's now do 10-fold cross validation to evaluate the models.

Hide

Hide

```
control = trainControl(method="repeatedcv", number=10, repeats=3)
model5 <- train(rings~., data = KNNtrain, method = "knn", preProcess="scale", trControl=control)
model5
```

## k-Nearest Neighbors

2901 samples

7 predictor

3 classes: 'young', 'adult', 'old'

Pre-processing: scaled (7)

Resampling: Cross-Validated (10 fold, repeated 3 times)

Summary of sample sizes: 2612, 2610, 2610, 2610, 2611, 2611, ...

Resampling results across tuning parameters:

k	Accuracy	Kappa
5	0.6564605	0.4565599
7	0.6689747	0.4743428
9	0.6789763	0.4889676

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was k = 9.

The 10-fold cross validation method indicates that the optimal model for KNN is one with k = 9 (same as what the bootstrap method predicted).

The 10-fold cross validation method indicates that the optimal model for Naive Bayes is a model with fL = 0, usekernal = TRUE and adjust = 1.

The cross-validation method confirms that the KNN method is more effective for this data set than Naive Bayes.

Let's create the new models with the suggested parameters.

Hide

Hide

```
library(class)
KNNpred <- knn(train = KNNtrain[1:7], test = KNNtest[1:7], cl = KNNtrain$strings, k = 9
)
library(caret)
confusionMatrix(KNNpred, KNNtest$strings)
```

Confusion Matrix and Statistics

Reference				
Prediction	young	adult	old	
young	322	96	18	
adult	90	375	164	
old	6	72	133	

Overall Statistics

Accuracy : 0.6505  
95% CI : (0.6236, 0.6767)  
No Information Rate : 0.4255  
P-Value [Acc > NIR] : < 2e-16  
  
Kappa : 0.4517  
Mcnemar's Test P-Value : 3.9e-09

Statistics by Class:

	Class: young	Class: adult	Class: old
Sensitivity	0.7703	0.6906	0.4222
Specificity	0.8671	0.6535	0.9188
Pos Pred Value	0.7385	0.5962	0.6303
Neg Pred Value	0.8857	0.7403	0.8291
Prevalence	0.3276	0.4255	0.2469
Detection Rate	0.2524	0.2939	0.1042
Detection Prevalence	0.3417	0.4929	0.1654
Balanced Accuracy	0.8187	0.6720	0.6705

With k = 9 the model was about 69% accurate in predicting the test data set.

Hide

Hide

```
library(e1071)
```

Hide

Hide

```
model <- naiveBayes(rings ~., data = NBtrain, fL = 0, usekernel = TRUE, adjust = 1)
model
```

Naive Bayes Classifier for Discrete Predictors

Call:  
naiveBayes.default(x = X, y = Y, laplace = laplace, fL = 0, usekernel = TRUE,

```
adjust = 1)
```

A-priori probabilities:

```
Y
    young    adult    old
0.3409169 0.4367459 0.2223371
```

Conditional probabilities:

length

```
Y          [,1]      [,2]
young 0.4716203 0.1490067
adult 0.6712654 0.1177351
old    0.6912424 0.1092433
```

diameter

```
Y          [,1]      [,2]
young 0.4507906 0.1502746
adult 0.6580688 0.1193139
old    0.6863657 0.1123703
```

height

```
Y          [,1]      [,2]
young 0.09493365 0.04008618
adult 0.13448953 0.02663204
old    0.14618234 0.02554316
```

whole\_weight

```
Y          [,1]      [,2]
young 0.1544802 0.1083561
adult 0.3494055 0.1509557
old    0.3924864 0.1602648
```

shucked\_wieght

```
Y          [,1]      [,2]
young 0.1337806 0.09876195
adult 0.2970106 0.14060814
old    0.2962226 0.13836111
```

viscera\_wieght

```
Y          [,1]      [,2]
young 0.1241098 0.09026025
adult 0.2866922 0.12776408
old    0.3145439 0.13646942
```

shell\_weight

```
Y          [,1]      [,2]
young 0.1206878 0.08019425
adult 0.2733141 0.10873391
old    0.3365036 0.13167884
```



```
pred <- predict(model, NBtest)
print(confusionMatrix(pred, NBtest$rings))
```

### Confusion Matrix and Statistics

```

      Reference
Prediction young adult old
  young    336    126  47
  adult     76    285 139
  old        6    132 129

```

### Overall Statistics

```

      Accuracy : 0.5878
      95% CI : (0.5602, 0.6149)
No Information Rate : 0.4255
P-Value [Acc > NIR] : < 2.2e-16

```

```

      Kappa : 0.3667
McNemar's Test P-Value : 1.32e-09

```

### Statistics by Class:

	Class: young	Class: adult	Class: old
Sensitivity	0.8038	0.5249	0.4095
Specificity	0.7984	0.7067	0.8564
Pos Pred Value	0.6601	0.5700	0.4831
Neg Pred Value	0.8931	0.6675	0.8157
Prevalence	0.3276	0.4255	0.2469
Detection Rate	0.2633	0.2234	0.1011
Detection Prevalence	0.3989	0.3918	0.2092
Balanced Accuracy	0.8011	0.6158	0.6330

With the suggested parameters given from the 10-fold validation, the naive bayes algorithm is about 59% accurate.

The models trained by the 10-fold validation have almost equal accuracy to the models I originally created, when testing on the test data set. My concern with this project is that the parameters I originally used didn't differ much from the suggested model in 10-fold validation.

```

control14 = trainControl(method="repeatedcv", number=10, repeats=3)
model7 <- train(rings~., data = KNNtrain, method = "rf", preProcess="scale", trControl=control14)

```

randomForest 4.6-12

Type `rfNews()` to see new features/changes/bug fixes.

Attaching package: 'randomForest'

The following object is masked from 'package:ggplot2':

margin

Hide

Hide

model7

Random Forest

2901 samples

7 predictor

3 classes: 'young', 'adult', 'old'

Pre-processing: scaled (7)

Resampling: Cross-Validated (10 fold, repeated 3 times)

Summary of sample sizes: 2611, 2611, 2611, 2612, 2611, 2610, ...

Resampling results across tuning parameters:

mtry	Accuracy	Kappa
2	0.6886180	0.5076597
4	0.6826480	0.4994915
7	0.6813857	0.4980003

Accuracy was used to select the optimal model using the largest value.

The final value used for the model was `mtry = 2`.

As shown above, the optimal random forest model is expected to perform at about 67%. It seems like the machine learning algorithms are having a difficult learning enough from the abalone features to accurately predict the age of the abalone. I would suspect that an abalone harvester would want a more accurate model before he or she could trust it in a commercial setting. Therefore, better data might be necessary.